


Chapter 2 Mutual Exclusion




Review: Amdahl's Law

$$\text{Speedup} = \frac{1}{1 - p + \frac{p}{n}}$$



© Instructor generated course materials (e.g., handouts, notes, summaries, exam questions, etc.) are protected by law and may not be copied or distributed in any form or in any medium without explicit permission of the instructor. Note that infringements of copyright can be subject to follow up by the University under the Code of Student Conduct and Disciplinary Procedures.

Shared Data Structures



Example Synchronization Paradigms

- Mutual exclusion
- Readers-Writers
- Producer-Consumer

Mutual Exclusion



- We will clarify our understanding of mutual exclusion
- We will also show how to reason about various properties in an asynchronous concurrent setting



Mutual Exclusion



In his 1965 paper E. W. Dijkstra wrote:

"Given in this paper is a solution to a problem which, to the knowledge of the author, has been an open question since at least 1962, irrespective of the solvability. [...] Although the setting of the problem might seem somewhat academic at first, the author trusts that anyone familiar with the logical problems that arise in computer coupling will appreciate the significance of the fact that this problem indeed can be solved."



Mutual Exclusion



- Formal problem definitions
- Solutions for 2 threads
- Solutions for n threads
- Fair solutions
- Inherent costs



Warning

- You will *never* use these protocols
 - Get over it
- You are advised to understand them
 - The same issues show up everywhere
 - Except hidden and more complex




Why is Concurrent Programming so Hard?

- Try preparing a seven-course banquet
 - By yourself
 - With one friend
 - With twenty-seven friends ...
- Before we can talk about programs
 - Need a language
 - Describing time and concurrency




Time

- *“Absolute, true and mathematical time, of itself and from its own nature, flows equably without relation to anything external.”* (Isaac Newton, 1689)
- *“Time is what keeps everything from happening at once.”* (Ray Cummings, 1922)




Events

- An *event* a_0 of thread A is
 - Instantaneous
 - No simultaneous events (break ties)



Threads

- A *thread* A is (formally) a sequence a_0, a_1, \dots of events
 - “Trace” model
 - Notation: $a_0 \rightarrow a_1$ indicates order




Example Thread Events

- Assign to shared variable
- Assign to local variable
- Invoke method
- Return from method
- Lots of other things ...



Threads are State Machines




States

- Thread State
 - Program counter
 - Local variables
- System state
 - Object fields (shared variables)
 - Union of thread states




Concurrency

- Thread A



Concurrency

- Thread A



- Thread B




Interleavings

- Events of two or more threads
 - Interleaved
 - Not necessarily independent (why?)




Intervals


- An *interval* $A_0 = (a_0, a_1)$ is
 - Time between events a_0 and a_1



Intervals may Overlap




Intervals may be Disjoint




Precedence

Interval A_0 precedes interval B_0




Precedence



- Notation: $A_0 \rightarrow B_0$
- Formally,
 - End event of A_0 before start event of B_0
 - Also called “happens before” or “precedes”




Precedence Ordering



- Remark: $A_0 \rightarrow B_0$ is just like saying
 - 1066 AD \rightarrow 1492 AD,
 - Middle Ages \rightarrow Renaissance,
- Oh wait,
 - what about this week vs this month?



Precedence Ordering



- Never true that $A \rightarrow A$
- If $A \rightarrow B$ then not true that $B \rightarrow A$
- If $A \rightarrow B$ & $B \rightarrow C$ then $A \rightarrow C$
- Funny thing: $A \rightarrow B$ & $B \rightarrow A$ might both be false!



Partial Orders

(review)

- Irreflexive:
 - Never true that $A \rightarrow A$
- Antisymmetric:
 - If $A \rightarrow B$ then not true that $B \rightarrow A$
- Transitive:
 - If $A \rightarrow B$ & $B \rightarrow C$ then $A \rightarrow C$



Total Orders

(review)

- Also
 - Irreflexive
 - Antisymmetric
 - Transitive
- Except that for every distinct A, B ,
 - Either $A \rightarrow B$ or $B \rightarrow A$



Repeated Events

```
while (mumble) {
  a0; a1;
}
```

k -th occurrence of
event a_0

a_0^k

k -th occurrence of
interval $A_0 = (a_0, a_1)$

A_0^k



Implementing a Counter

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        temp = value;  
        value = temp + 1;  
        return temp;  
    }  
}
```

Make these steps
indivisible using locks



Locks (Mutual Exclusion)

```
public interface Lock {  
  
    public void lock();  
  
    public void unlock();  
}
```



Locks (Mutual Exclusion)

```
public interface Lock {  
    public void lock();  
    public void unlock();  
}
```

acquire lock



Locks (Mutual Exclusion)

```
public interface Lock {  
    public void lock();  
    public void unlock();  
}
```

acquire lock

release lock



Using Locks

```
public class Counter {
    private long value;
    private Lock lock;    //to protect critical section
    public long getAndIncrement() {
        lock.lock();      //enter critical section
        try {
            int temp = value;    //in critical section
            value = value + 1;
        } finally {
            lock.unlock();      //leave critical section
        }
        return temp;
    }
}
```



Using Locks

```
public class Counter {
    private long value;
    private Lock lock;
    public long getAndIncrement() {
        lock.lock();
        try {
            int temp = value;
            value = value + 1;
        } finally {
            lock.unlock();
        }
        return temp;
    }
}
```

acquire Lock



Using Locks

```
public class Counter {
    private long value;
    private Lock lock;
    public long getAndIncrement() {
        lock.lock();
        try {
            int temp = value;
            value = value + 1;
        } finally {
            lock.unlock();
        }
        return temp;
    }
}
```

Release lock
(no matter what)




Using Locks

```
public class Counter {
    private long value;
    private Lock lock;
    public long getAndIncrement() {
        lock.lock();
        try {
            int temp = value;
            value = value + 1;
        } finally {
            lock.unlock();
        }
        return temp;
    }
}
```

critical section





Mutual Exclusion

- Let CS_i^k  be thread i's k-th critical section execution









Mutual Exclusion

- Let CS_i^k  be thread i's k-th critical section execution
- And CS_j^m  be thread j's m-th critical section execution









Mutual Exclusion

- Let CS_i^k  be thread i's k-th critical section execution
- And CS_j^m  be j's m-th execution
- Then either
 -   or  





Mutual Exclusion

- Let CS_i^k  be thread i's k-th critical section execution
- And CS_j^m  be j's m-th execution
- Then either
 -   or  

$CS_i^k \rightarrow CS_j^m$



Mutual Exclusion

- Let CS_i^k  be thread i's k-th critical section execution
- And CS_j^m  be j's m-th execution
- Then either

—   or  

$CS_i^k \rightarrow CS_j^m$

$CS_j^m \rightarrow CS_i^k$



Deadlock-Free



- If some thread calls `lock()`
 - And never returns
 - Then other threads must complete `lock()` and `unlock()` calls infinitely often
- System as a whole makes progress
 - Even if individuals starve



Starvation-Free



- If some thread calls `lock()`
 - It will eventually return
- Individual threads make progress



Two-Thread vs n -Thread Solutions

- 2-thread solutions first
 - Illustrate most basic ideas
 - Fits on one slide
- Then n -thread solutions



Two-Thread Conventions

```
class ... implements Lock {  
    ...  
    // thread-local index, 0 or 1  
    public void lock() {  
        int i = ThreadID.get();  
        int j = 1 - i;  
        ...  
    }  
}
```



Two-Thread Conventions

```
class ... implements Lock {  
    ...  
    // thread-local index, 0 or 1  
    public void lock() {  
        int i = ThreadID.get();  
        int j = 1 - i;  
        ...  
    }  
}
```

Henceforth: *i* is current
thread, *j* is other thread



LockOne

```
class LockOne implements Lock {  
    private boolean[] flag = new boolean[2];  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```



LockOne

```
class LockOne implements Lock {  
    private boolean[] flag = new boolean[2];  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```

Each thread has flag



LockOne

```
class LockOne implements Lock {  
    private boolean[] flag = new boolean[2];  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```

Set my flag



LockOne

```
class LockOne implements Lock {  
    private boolean[] flag = new boolean[2];  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```

**Wait for other flag to
become false**



LockOne Satisfies Mutual Exclusion

- Assume CS_A^j overlaps CS_B^k
- Consider each thread's last
 - (j^{th} and k^{th}) read and write ...
 - in `lock()` before entering
- Derive a contradiction



From the Code

- $\text{write}_A(\text{flag}[A]=\text{true}) \rightarrow$
 $\text{read}_A(\text{flag}[B]==\text{false}) \rightarrow CS_A$
- $\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow$
 $\text{read}_B(\text{flag}[A]==\text{false}) \rightarrow CS_B$

```
class LockOne implements Lock {
    ...
    public void lock() {
        flag[i] = true;
        while (flag[j]) {}
    }
}
```



From the Assumption

- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$

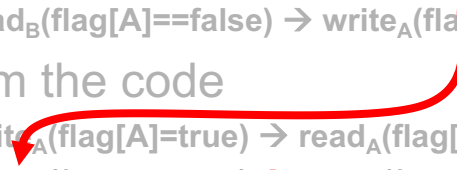


Combining

- Assumptions:
 - $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
 - $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$
- From the code
 - $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
 - $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$

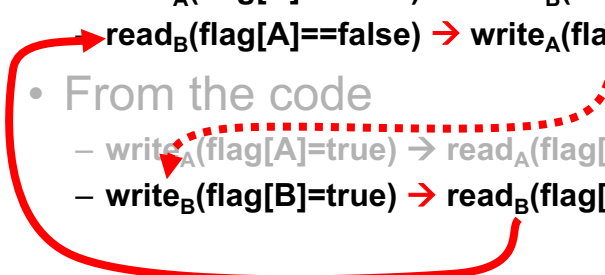


Combining

- Assumptions:
 - $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
 - $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$
 - From the code
 - $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
 - $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$
- 



Combining

- Assumptions:
 - $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
 - $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$
 - From the code
 - $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
 - $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$
- 



Combining

- Assumptions:

- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$

- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$

- From the code

- $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$

- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$



Combining

- Assumptions:

- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$

- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$

- From the code

- $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$

- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$



Combining

- Assumptions.

- ~~read_A(flag[B]==false)~~ → write_B(flag[B]=true)

- ~~read_B(flag[A]==false)~~ → write_A(flag[A]=true)


- From the code

- ~~write_A(flag[A]=true)~~ → read_A(flag[B]==false)

- ~~write_B(flag[B]=true)~~ → read_B(flag[A]==false)



Cycle!



Impossible in a
partial order



LockTwo

```
public class LockTwo implements Lock {  
    private int victim;  
    public void lock() {  
        victim = i;  
        while (victim == i) {};  
    }  
  
    public void unlock() {}  
}
```

Let other go first



LockTwo

```
public class LockTwo implements Lock {  
    private int victim;  
    public void lock() {  
        victim = i;  
        while (victim == i) {};  
    }  
  
    public void unlock() {}  
}
```

Wait for permission



LockTwo

```
public class Lock2 implements Lock {
    private int victim;
    public void lock() {
        victim = i;
        while (victim == i) {};
    }
    public void unlock() {}
}
```

Nothing to do



LockTwo Claims

- Satisfies mutual exclusion

- If thread *i* in CS
- Then `victim == i`
- Cannot be both 0 and 1

```
public void LockTwo() {
    victim = i;
    while (victim == i) {};
}
```

- Not deadlock free

- Sequential execution deadlocks
- Concurrent execution does not



Peterson's Algorithm

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
public void unlock() {  
    flag[i] = false;  
}
```



Peterson's Algorithm

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
public void unlock() {  
    flag[i] = false;  
}
```

**Announce I'm
interested**



Peterson's Algorithm

```

public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == i) {};
}
public void unlock() {
    flag[i] = false;
}

```

Announce I'm interested (points to `flag[i] = true;`)

Defer to other (points to `victim = i;`)



Peterson's Algorithm

```

public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == i) {};
}
public void unlock() {
    flag[i] = false;
}

```

Announce I'm interested (points to `flag[i] = true;`)

Defer to other (points to `victim = i;`)

Wait while other interested & I'm the victim (points to `while (flag[j] && victim == i) {};`)



Peterson's Algorithm

```

public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == i) {};
}
public void unlock() {
    flag[i] = false;
}

```

Announce I'm interested (points to `flag[i] = true;`)

Defer to other (points to `while (flag[j] && victim == i) {};`)

Wait while other interested & I'm the victim (points to `while (flag[j] && victim == i) {};`)

No longer interested (points to `flag[i] = false;`)



Mutual Exclusion

(1) $\text{write}_B(\text{Flag}[B]=\text{true}) \rightarrow \text{write}_B(\text{victim}=B)$

```

public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == i) {};
}

```

From the Code



Also from the Code

(2) $\text{write}_A(\text{victim}=A) \rightarrow \text{read}_A(\text{flag}[B])$
 $\rightarrow \text{read}_A(\text{victim})$

```
public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == i) {};
}
```



Assumption

(3) $\text{write}_B(\text{victim}=B) \rightarrow \text{write}_A(\text{victim}=A)$

W.L.O.G. assume **A** is the last
 thread to write **victim**



Combining Observations

- (1) $\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow \text{write}_B(\text{victim}=B)$
- (3) $\text{write}_B(\text{victim}=B) \rightarrow \text{write}_A(\text{victim}=A)$
- (2) $\text{write}_A(\text{victim}=A) \rightarrow \text{read}_A(\text{flag}[B])$
 $\rightarrow \text{read}_A(\text{victim})$



Combining Observations

- (1) $\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow$
- (3) $\text{write}_B(\text{victim}=B) \rightarrow$
- (2) $\text{write}_A(\text{victim}=A) \rightarrow \text{read}_A(\text{flag}[B])$
 $\rightarrow \text{read}_A(\text{victim})$



Combining Observations

(1) $\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow$

(3) $\text{write}_B(\text{victim}=B) \rightarrow$

(2) $\text{write}_A(\text{victim}=A) \rightarrow \text{read}_A(\text{flag}[B])$
 $\rightarrow \text{read}_A(\text{victim})$

A read $\text{flag}[B] == \text{true}$ and $\text{victim} == A$, so it could not have entered the CS (QED)



Deadlock Free

```
public void lock() {
    ...
    while (flag[j] && victim == i) {};
```

- Thread blocked
 - only at **while** loop
 - only if other's flag is true
 - only if it is the victim
- Solo: other's flag is false
- Both: one or the other not the victim



Starvation Free

- Thread i blocked only if j repeatedly re-enters so that

$\text{flag}[j] == \text{true}$ and $\text{victim} == i$

- When j re-enters
 - it sets victim to j .
 - So i gets in

```
public void lock() {
    flag[i] = true;
    victim  = i;
    while (flag[j] && victim == i) {};
}


public void unlock() {
    flag[i] = false;
}
```



The Filter Algorithm for n Threads

There are $n-1$ “waiting rooms” called levels

- At each level
 - At least one enters level
 - At least one blocked if many try
- Only one thread makes it through



Filter

```

class Filter implements Lock {
    int[] level; // level[i] for thread i
    int[] victim; // victim[L] for level L

    public Filter(int n) {
        level = new int[n];
        victim = new int[n];
        for (int i = 1; i < n; i++) {
            level[i] = 0;
        }
        ...
    }
}

```

Thread 2 at level 4



Filter

```

class Filter implements Lock {
    ...

    public void lock() {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while ((∃ k != i level[k] >= L) &&
                    victim[L] == i) {};
        }
    }

    public void unlock() {
        level[i] = 0;
    }
}

```



Filter

```

class Filter implements Lock {
    ...

    public void lock() {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while (( $\exists$  k != i) level[k] >= L) &&
                victim[L] == i) {}
        }
    }
    public void release(int i) {
        level[i] = 0;
    }
}

```

One level at a time



Filter

```

class Filter implements Lock {
    ...

    public void lock() {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while (( $\exists$  k != i) level[k] >= L) &&
                victim[L] == i) {}
        }
    }
    public void release(int i) {
        level[i] = 0;
    }
}

```

Announce intention to enter level L



Filter

```

class Filter implements Lock {
    int level[n];
    int victim[n];
    public void lock() {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while (( $\exists$  k != i) level[k] >= L) &&
                victim[L] == i) {};
        }
    }
    public void release(int i) {
        level[i] = 0;
    }
}

```

Give priority to anyone but me



Filter

Wait as long as someone else is at same or higher level, and I'm designated victim

```

public void lock() {
    for (int L = 1; L < n; L++) {
        level[i] = L;
        victim[L] = i;
        while (( $\exists$  k != i) level[k] >= L) &&
            victim[L] == i) {};
    }
}
public void release(int i) {
    level[i] = 0;
}
}

```



Filter


```
class Filter implements Lock {
    int level[n];
    int victim[n];
    public void lock() {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while (( $\exists$  k != i) level[k] >= L) &&
                victim[L] == i) {};
        }
    }
}
```

Thread enters level L when it completes the loop




Claim

- Start at level $L=0$
- At most $n-L$ threads enter level L
- Mutual exclusion at level $L=n-1$




Induction Hypothesis

- No more than $n-(L-1)$ at level L-1
- Induction step: by contradiction
- Assume all at level L-1 enter level L
- A last to write `victim[L]`
- B is any other thread at level L



Proof Structure



Show that A must have seen B in `level[L]` and since `victim[L] == A` could not have entered



Just Like Peterson

(1) $\text{write}_B(\text{level}[B]=L) \rightarrow \text{write}_B(\text{victim}[L]=B)$

```
public void lock() {
    for (int L = 1; L < n; L++) {
        level[i] = L;
        victim[L] = i;
        while (( $\exists k \neq i$ ) level[k] >= L
            && victim[L] == i) {};
    }
}
```

From the Code



From the Code

(2) $\text{write}_A(\text{victim}[L]=A) \rightarrow \text{read}_A(\text{level}[B])$
 $\rightarrow \text{read}_A(\text{victim}[L])$

```
public void lock() {
    for (int L = 1; L < n; L++) {
        level[i] = L;
        victim[L] = i;
        while (( $\exists k \neq i$ ) level[k] >= L
            && victim[L] == i) {};
    }
}
```



By Assumption

(3) $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$

By assumption, A is the last thread
to write **victim[L]**



Combining Observations

(1) $\text{write}_B(\text{level}[B]=L) \rightarrow \text{write}_B(\text{victim}[L]=B)$

(3) $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$

(2) $\text{write}_A(\text{victim}[L]=A) \rightarrow \text{read}_A(\text{level}[B])$
 $\rightarrow \text{read}_A(\text{victim}[L])$



Combining Observations

- (1) $\text{write}_B(\text{level}[B]=L) \rightarrow$
 (3) $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$
 (2) $\rightarrow \text{read}_A(\text{level}[B])$
 $\rightarrow \text{read}_A(\text{victim}[L])$



Combining Observations

- (1) $\text{write}_B(\text{level}[B]=L) \rightarrow$
 (3) $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$
 (2) $\rightarrow \text{read}_A(\text{level}[B])$
 $\rightarrow \text{read}_A(\text{victim}[L])$

A read $\text{level}[B] \geq L$, and $\text{victim}[L] = A$, so it could not have entered level L!



No Starvation

- Filter Lock satisfies properties:
 - Just like Peterson Alg at any level
 - So no one starves
- But what about fairness?
 - Threads can be overtaken by others



Bounded Waiting

- Want stronger fairness guarantees
- Thread not “overtaken” too much
- If A starts before B, then A enters before B?
- But what does “start” mean?
- Need to adjust definitions



Bounded Waiting

- Divide `lock()` method into 2 parts:
 - Doorway interval:
 - Written D_A
 - always finishes in finite steps
 - Waiting interval:
 - Written W_A
 - may take unbounded steps



r-Bounded Waiting

- For threads A and B:
 - If $D_A^k \rightarrow D_B^j$
 - A's k-th doorway precedes B's j-th doorway
 - Then $CS_A^k \rightarrow CS_B^{j+r}$
 - A's k-th critical section precedes B's j+r-th critical section
 - B cannot overtake A more than r times
- First-come-first-served $\rightarrow r = 0$



What is “r” for Peterson's Algorithm?

```
public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == i) {};
}
public void unlock() {
    flag[i] = false;
}
```

Answer: $r = 0$



What is “r” for the Filter Algorithm?

```
class Filter implements Lock {
    ...

    public void lock() {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while ((∃ k != i level[k] >= L) &&
                    victim[L] == i) {};
        }
    }
    public void unlock() {
        level[i] = 0;
    }
}
```

Answer: there is no value of “r”



First-Come-First-Served

- For threads A and B:
 - If $D_A^k \rightarrow D_B^j$
 - A's k-th doorway precedes B's j-th doorway
 - Then $CS_A^k \rightarrow CS_B^j$
 - A's k-th critical section precedes B's j-th critical section
 - B cannot overtake A



Fairness

- Filter Lock satisfies properties:
 - No one starves
 - But very weak fairness
 - Can be overtaken **arbitrary** # of times
 - So being fair is stronger than avoiding starvation
 - And filter is pretty lame...



Bakery Algorithm

- Provides First-Come-First-Served for n threads
- How?
 - Take a “number”
 - Wait until lower numbers have been served
- Lexicographic order
 - $(a,i) > (b,j)$
 - If $a > b$, or $a = b$ and $i > j$



Bakery Algorithm

```
class Bakery implements Lock {
    boolean[] flag;
    Label[] label;
    public Bakery (int n) {
        flag = new boolean[n];
        label = new Label[n];
        for (int i = 0; i < n; i++) {
            flag[i] = false; label[i] = 0;
        }
    }
    ...
}
```



Bakery Algorithm

```
class Bakery implements Lock {
    boolean[] flag;
    Label[] label;
    public Bakery (int n) {
        flag = new boolean[n];
        label = new Label[n];
        for (int i = 0; i < n; i++) {
            flag[i] = false; label[i] = 0;
        }
        ...
    }
}
```

Diagram illustrating the Bakery Algorithm state. Two processors, labeled 2 (red) and 6 (blue), are shown. Below them is a row of 8 boxes representing the flag array: [f, f, t, f, f, t, f, f]. Below that is a row of 8 boxes representing the label array: [0, 0, 4, 0, 0, 5, 0, 0]. Arrows point from the 't' in the flag array and the '4' and '5' in the label array to a central 'CS' (Critical Section) label.



Bakery Algorithm

```
class Bakery implements Lock {
    ...
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1]) + 1;
        while (∃k flag[k]
            && (label[i], i) > (label[k], k));
    }
}
```



Bakery Algorithm

```
class Bakery implements Lock {
    ...
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1])+1;
        while (∃k flag[k]
                && (label[i],i) > (label[k],k));
    }
}
```

Doorway



Bakery Algorithm

```
class Bakery implements Lock {
    ...
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1])+1;
        while (∃k flag[k]
                && (label[i],i) > (label[k],k));
    }
}
```

I'm interested



Bakery Algorithm

Take increasing
label (read labels
in some arbitrary
order)

```
class Bakery implements Lock {
    ...
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1])+1;
        while (∃k flag[k]
                && (label[i], i) > (label[k], k));
    }
}
```



Bakery Algorithm

Someone is
interested

```
class Bakery implements Lock {
    ...
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1])+1;
        while (∃k flag[k]
                && (label[i], i) > (label[k], k));
    }
}
```



Bakery Algorithm

```
class Bakery implements Lock {
    boolean flag[n];
    int label[n];

    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1])+1;
        while ( $\exists k$  flag[k]
            && (label[i],i) > (label[k],k));
    }
}
```

Someone is
interested ...

... whose (label,i) in
lexicographic order is lower



Bakery Algorithm

```
class Bakery implements Lock {

    ...

    public void unlock() {
        flag[i] = false;
    }
}
```



Bakery Algorithm

```
class Bakery implements Lock {  
    ...  
    public void unlock() {  
        flag[i] = false;  
    }  
}
```

No longer interested

labels are always increasing



No Deadlock

- There is always one thread with earliest label
- Ties are impossible (why?)



First-Come-First-Served

- If $D_A \rightarrow D_B$ then
 - A's label is smaller
- And:
 - $\text{write}_A(\text{label}[A]) \rightarrow$
 - $\text{read}_B(\text{label}[A]) \rightarrow$
 - $\text{write}_B(\text{label}[B]) \rightarrow \text{read}_B(\text{flag}[A])$
- So B sees
 - smaller label for A
 - locked out while $\text{flag}[A]$ is true

```
class Bakery implements Lock {
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0],
            ..., label[n-1]) + 1;
        while (exists k flag[k]
            && (label[i], i) >
            (label[k], k));
    }
}
```



Mutual Exclusion

- Suppose A and B in CS together
- Suppose A has earlier label
- When B entered, it must have seen
 - $\text{flag}[A]$ is false, or
 - $\text{label}[A] > \text{label}[B]$

```
class Bakery implements Lock {
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0],
            ..., label[n-1]) + 1;
        while (exists k flag[k]
            && (label[i], i) >
            (label[k], k));
    }
}
```



Mutual Exclusion

- Labels are strictly increasing so
- B must have seen $\text{flag}[A] == \text{false}$



Mutual Exclusion

- Labels are strictly increasing so
- B must have seen $\text{flag}[A] == \text{false}$
- $\text{Labeling}_B \rightarrow \text{read}_B(\text{flag}[A]) \rightarrow \text{write}_A(\text{flag}[A]) \rightarrow \text{Labeling}_A$



Mutual Exclusion

- Labels are strictly increasing so
- B must have seen $\text{flag}[A] == \text{false}$
- $\text{Labeling}_B \rightarrow \text{read}_B(\text{flag}[A]) \rightarrow \text{write}_A(\text{flag}[A]) \rightarrow \text{Labeling}_A$
- Which contradicts the assumption that A has an earlier label



Bakery Y2³²K Bug

```
class Bakery implements Lock {
    ...
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1])+1;
        while (∃k flag[k]
                && (label[i], i) > (label[k], k));
    }
}
```



Bakery Y2³²K Bug

```
class Bakery implements Lock {
    ...
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1]) + 1;
        while (∃k flag[k]
                && (label[i], i) > (label[k], k));
    }
}
```

Mutex breaks if label[i] overflows?



Does Overflow Actually Matter?

- Yes
 - Y2K
 - 18 January 2038 (Unix `time_t` rollover)
 - 16-bit counters
- No
 - 64-bit counters
- Maybe
 - 32-bit counters



Timestamps

- Label variable is really a **timestamp**
- Need ability to
 - Read others' timestamps
 - Compare them
 - Generate a **later** timestamp
- Can we do this without overflow?



The Good News

- One can construct a
 - Wait-free (no mutual exclusion)
 - Concurrent
 - Timestamping system
 - That never overflows



The ~~Good~~ News

- One can construct a
 - Wait-free (no mutual exclusion)
 - Concurrent
 - Timestamping system
 - That never overflows
- This part is hard




Instead ...

- We construct a Sequential timestamping system
 - Same basic idea
 - But simpler
- As if we use mutex to read & write atomically
- No good for building locks
 - But useful anyway




Precedence Graphs



- Timestamps form directed graph
- Edge x to y
 - Means x is later timestamp
 - We say x **dominates** y




Unbounded Counter Precedence Graph




- Timestamping = move tokens on graph
- Atomically
 - read others' tokens
 - move mine
- Ignore tie-breaking for now




Unbounded Counter Precedence Graph




Unbounded Counter Precedence Graph




Two-Thread Bounded Precedence Graph




Two-Thread Bounded Precedence Graph




Two-Thread Bounded Precedence Graph




Two-Thread Bounded Precedence Graph




Two-Thread Bounded Precedence Graph T^2




Three-Thread Bounded Precedence Graph?




Three-Thread Bounded Precedence Graph?




Graph Composition




Three-Thread Bounded Precedence Graph T^3



Three-Thread Bounded Precedence Graph T^3




Three-Thread Bounded Precedence Graph T^3



In General

$T^k = T^2 * T^{k-1}$
K threads need 3^k nodes

label size =
 $\text{Log}_2(3^k) = 2k$



Deep Philosophical Question

- The Bakery Algorithm is
 - Succinct,
 - Elegant, and
 - Fair.
- Q: So why isn't it practical?
- A: Well, you have to read N distinct variables



Shared Memory

- Shared read/write memory locations called *Registers* (historical reasons)
- Come in different flavors
 - Multi-Reader-Single-Writer (`flag[]`)
 - Multi-Reader-Multi-Writer (`victim[]`)
 - Not that interesting: SRMW and SRSW



Theorem

At least N MRSW (multi-reader/single-writer) registers are needed to solve deadlock-free mutual exclusion.

N registers such as `flag[]`...



Proving Algorithmic Impossibility

- To show no algorithm exists:
 - assume by way of contradiction one does,
 - show a **bad execution** that violates properties:
 - in our case assume an alg for deadlock free mutual exclusion using $< N$ registers




CS



Proof: Need N-MRSW Registers

Each thread must write to some register



...can't tell whether **A** is in critical section



Upper Bound

- Bakery algorithm
 - Uses **2N** MRSW registers
- So the bound is (pretty) tight
- But what if we use MRMW registers?
 - Like `victim[]` ?



Bad News Theorem

At least N MRMW multi-reader/**multi-writer** registers are needed to solve deadlock-free mutual exclusion.

(So multiple writers don't help)




Theorem (For 2 Threads)

Theorem: **Deadlock-free mutual exclusion for 2 threads requires at least 2 multi-reader multi-writer registers**

Proof: **assume one register suffices and derive a contradiction**




Two Thread Execution



- Threads run, reading and writing R
- Deadlock free so at least one gets in



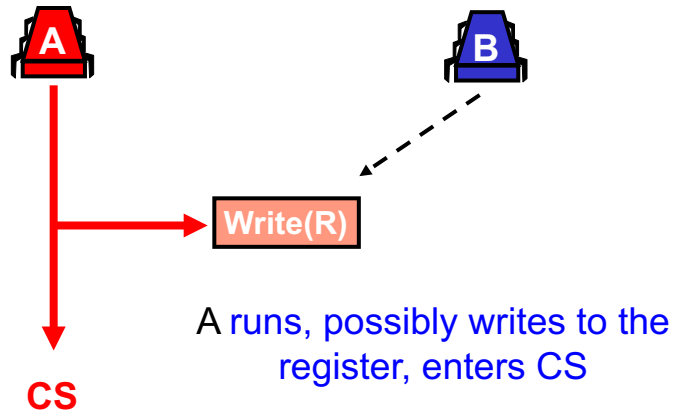
Covering State for One Register Always Exists



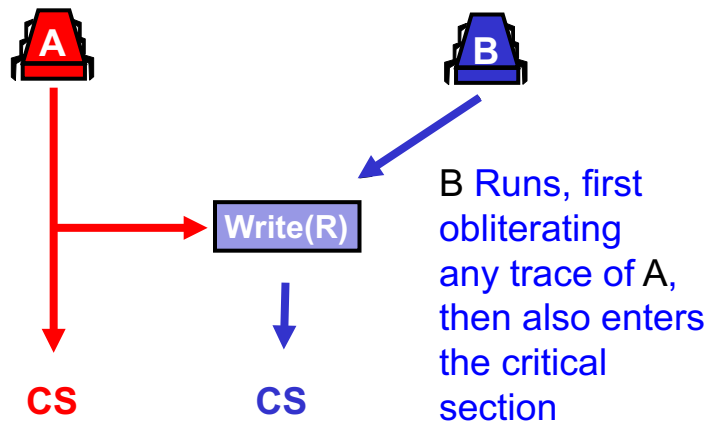
In any protocol B has to write to the register before entering CS, so stop it just before



Proof: Assume Cover of 1



Proof: Assume Cover of 1



Theorem

Deadlock-free mutual exclusion for 3 threads requires at least 3 multi-reader multi-writer registers

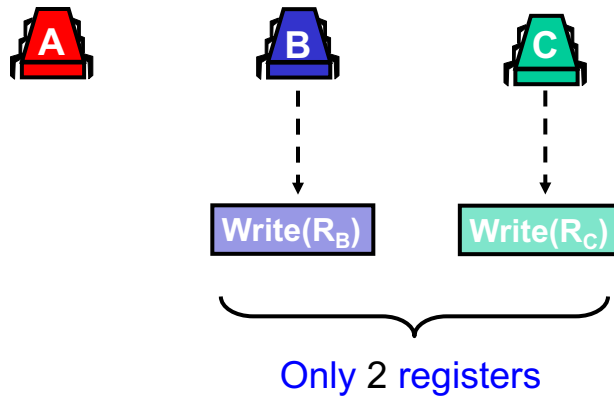


Theorem

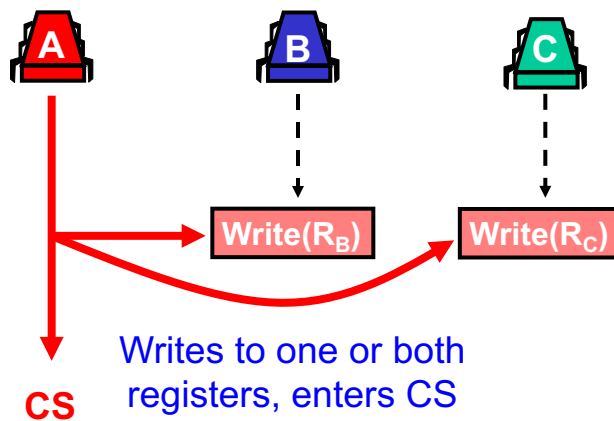
Deadlock-free mutual exclusion for n threads requires at least n multi-reader multi-writer registers




Proof: Assume Cover of 2



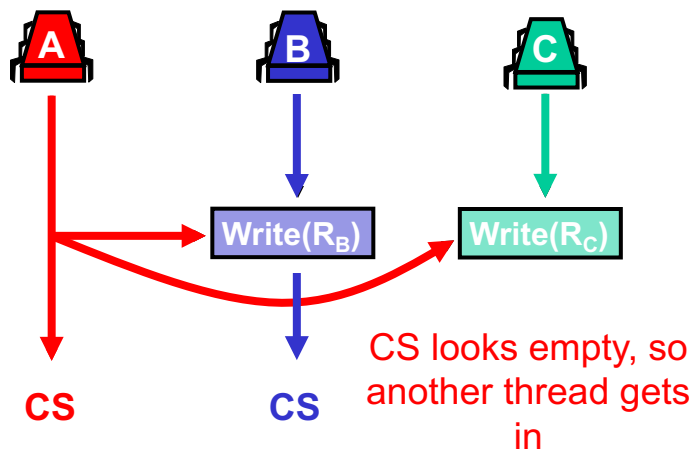
Run A Solo



Obliterate Traces of A



Mutual Exclusion Fails




Proof Strategy

- Proved: a contradiction starting from a covering state for 2 registers
- Claim: a covering state for 2 registers is reachable from any state where CS is empty




Covering State for Two



- If we run B through CS 3 times, B must return twice to cover some register, say R_B




Covering State for Two



- Start with B covering register R_B for the 1st time
- Run A until it is about to write to uncovered R_A
- Are we done?




Covering State for Two



- **NO!** A could have written to R_B
- So CS no longer looks empty to thread C




Covering State for Two



- Run **B** obliterating traces of **A** in R_B
- Run **B** again until it is about to write to R_B
- Now we are done



Inductively We Can Show



- There is a covering state
 - Where k threads not in CS cover k distinct registers
 - Proof follows when $k = N-1$



Summary of Lecture

- In the 1960's several **incorrect** solutions to starvation-free mutual exclusion using RW-registers were published...
- Today we know how to solve FIFO N thread mutual exclusion using $2N$ RW-Registers



Summary of Lecture

- N RW-Registers inefficient
 - Because writes **“cover”** older writes
- Need stronger hardware operations
 - that do not have the **“covering problem”**
- In next lectures - understand what these operations are...

